

Nvidia GPU Computation

COMP4901Y

Binhang Yuan

GPU Architecture

ANNOUNCING NVIDIA A100 PCIE

Greatest Generational Leap - 20X Volta


	Peak		Vs Volta
FP32 TRAINING	312 TFLOPS		20X
INT8 INFERENCE	1,248 TOPS		20X
FP64 HPC	19.5 TFLOPS		2.5X
MULTI INSTANCE GPU			7X GPUs




54B XTOR | 826mm² | TSMC 7N | 40GB Samsung HBM2

GPU Architecture


- The GPU is a highly parallel processor architecture, including processing elements and a memory hierarchy.
- The memory hierarchy:
 - L0, L1 cache in Streaming Multiprocessors (SMs);
 - On-chip L2 cache;
 - High bandwidth DRAM (HBM).
- Arithmetic and other instructions are executed by the SMs.
- Data and code are accessed from DRAM via the L2 cache.



Ampere GPU Architecture



108 SM in a A100 GPU



Ampere GPU SM



- In Ampere GPU, SM contains **four** processing blocks that share an L1 cache for data caching.
- Each processing block has:
 - 1 Warp scheduler (where the maximum number of thread blocks per SM is 32);
 - 16 INT32 CUDA cores;
 - 16 FP32 CUDA cores;
 - 8 FP64 CUDA cores;
 - 8 Load/Store cores;
 - 1 SFU core (special function units: e.g., sin, cos)
 - 1 **Tensor core** for matrix multiplication;
 - 1 16K 32-bit register file.




A100 GPU Memory Hierarchy

- Size:
 - L1 cache: 192 KB per SM;
 - L2 cache: 40 MB
 - HBM: 80 GB
- Accessibility:
 - The L2 cache is unified, shared by all SMs, and set aside for data and instructions.
 - The L1 instruction cache is private to a single streaming multiprocessor.
 - The L0 instruction cache is private to a single streaming multiprocessor subprocessing block.

<https://images.nvidia.com/aem-dam/en-zz/Solutions/data-center/nvidia-ampere-architecture-whitepaper.pdf>

A100 GPU Tensor Core Computation

- Multiply-add is the most frequent operation in modern neural networks. This is known as the fused multiply-add (FMA) operation.
- Includes one multiply operation and one add operation, counted as two float operations.
- A100 GPU has 1.41 GHz clock rate.
- The Ampere A100 GPU Tensor Cores multiply-add operations per clock:



Ampere A100 GPU FMA per clock on a SM					
FP64	TF32	FP16	INT8	INT4	INT1
64	512	1024	2048	4096	16384

A100 GPU Specifications

Ampere A100 GPU FMA per clock on a SM					
FP64	TF32	FP16	INT8	INT4	INT1
64	512	1024	2048	4096	16384

A100 GPU Specs	
Tensor core Float 32 (TF32)	156 TFLOPS
Tensor core Float 16 (FP16)	312 TFLOPS
Tensor core Int 8 (INT8)	624 TOPS
GPU Memory	80 GB
GPU Memory Bandwidth	2039 GB/s

$1024 \times 2 \times 1.41 \times 10^9 \times 108 = 312 \times 10^{12}$


Tensor Cores

- Tensor Cores were introduced in the NVIDIA Volta™ GPU architecture to accelerate matrix multiply and accumulate operations for machine learning and scientific applications.
- These instructions operate on small matrix blocks:
 - For example, 16×16 blocks in A100 GPUs.
- Tensor Cores can compute and accumulate products with higher precision than the inputs:
 - During training with FP16 inputs, Tensor Cores can compute products without loss of precision and accumulate in FP32.

H100 GPU

SPECIFICATIONS



	H100 SXM	H100 PCIe
FP64	34 TFLOPS	26 TFLOPS
FP64 Tensor Core	67 TFLOPS	51 TFLOPS
FP32	67 TFLOPS	51 TFLOPS
TF32 Tensor Core	989 TFLOPS*	756 TFLOPS*
BFLOAT16 Tensor Core	1,979 TFLOPS*	1,513 TFLOPS*
FP16 Tensor Core	1,979 TFLOPS*	1,513 TFLOPS*
FP8 Tensor Core	3,958 TFLOPS*	3,026 TFLOPS*
INT8 Tensor Core	3,958 TOPS*	3,026 TOPS*
GPU memory	80GB	80GB
GPU memory bandwidth	3.35TB/s	2TB/s




* Shown with sparsity. Specifications 1/2 lower without sparsity.



H100 GPU - Highlights


- New fourth-generation Tensor Cores are up to 6x faster chip-to-chip compared to A100 (the third-generation), including per-SM speedup, additional SM count, and higher clocks of H100.
- Hopper FP8 Data Format:
 - Add FP8 Tensor Cores to accelerate both AI training and inference;
 - Two new FP8 inputs types:
 - E4M3 with 4 exponent bits, 3 mantissa bits, and 1 sign bit
 - E5M2 with 5 exponent bits, 2 mantissa bits, and 1 sign bit.
 - Support multiple accumulator and output types.










What is next? GB200 soon in 2025...

 **NVIDIA** >

Shop Drivers Support  


Explore what's next in AI and accelerated computing at GTC, March 17–21.
See recommended sessions. 


Register Now 

Cloud & Data Center Solutions  Products  Data Center GPUs  Software  Technologies  Resources  Get Started

NVIDIA GB200 NVL72

Powering the new era of computing.

Read Datasheet 



Technical Specifications¹	
GB200 NVL72	
Blackwell GPUs Grace CPUs	72 36
CPU Cores	2,592 Arm Neoverse V2 Cores
Total FP4 Tensor Core	1,440 petaFLOPS
Total FP8/FP6 Tensor Core	720 petaFLOPS/petaOPS
Total Fast Memory	Up to 30TB
Total Memory Bandwidth	Up to 576TB/s
Total NVLink Bandwidth	130TB/s
Individual Blackwell GPU Specifications	
FP4 Tensor Core	20 petaFLOPS
FP8/FP6 Tensor Core	10 petaFLOPS
INT8 Tensor Core	10 petaOPS
FP16/BF16 Tensor Core	5 petaFLOPS
TF32 Tensor Core	2.5 petaFLOPS
FP32	80 teraFLOPS
FP64/FP64 Tensor Core	40 teraFLOPS
GPU Memory Bandwidth	186GB HBM3e 8 TB/s

Get Started

Stay Up to Date

Sign up to hear when NVIDIA Blackwell becomes available.

Notify Me


GPU Execution Model

2-level Thread Hierarchy

- GPUs execute functions using a 2-level hierarchy of threads.
 - A given function's threads are grouped into equally sized thread blocks, and a set of thread blocks is launched to execute the function.
- GPUs hide dependent instruction latency by switching to the execution of other threads.
 - The number of threads needed to utilize a GPU effectively is much higher than the number of cores or instruction pipelines.

Host: launch the function

Device: Parallel Kernel




2-level Thread Hierarchy

- GPUs have many SMs, each of which has pipelines for executing many threads and enables its threads to communicate via shared memory and synchronization.
- At runtime, a thread block is placed on an SM for execution, enabling all threads in a thread block to communicate and synchronize efficiently.
- Launching a function with a single thread block would only give work to a single SM; to fully utilize a GPU with multiple SMs, one needs to launch many thread blocks.
- Since an SM can execute multiple thread blocks concurrently, typically, one wants the number of thread blocks to be several times higher than the number of SMs.

2-level Thread Hierarchy

- Minimize the “tail” effect: at the end of a function execution, only a few active thread blocks remain.
- We use the term wave to refer to a set of thread blocks that run concurrently.
- It is most efficient to launch functions that execute in several waves of thread blocks - a smaller percentage of time is spent in the tail wave, minimizing the tail effect and thus the need to do anything about it.
- For the higher-end GPUs, typically only launches with fewer than 300 thread blocks should be examined for tail effects.



Utilization of an 8-SM GPU when 12 thread blocks with an occupancy of 1 block/SM at a time are launched for execution. The blocks execute in 2 waves, the first wave utilizes 100% of the GPU, while the 2nd wave utilizes only 50%.

Understanding Performance

Overview

- The performance of a function on a given processor is limited by one of the following three factors:
 - Memory bandwidth;
 - Math bandwidth;
 - Latency.
- Consider a simplified model where a function:
 - Read its input from memory;
 - Perform math operations;
 - Write its output to memory.

Modeling the Cost

- T_{mem} time is spent in accessing memory;
- T_{math} time is spent performing math operations.
- If we further assume that memory and math portions of different threads can be overlapped;
- The total time for the function is $\max(T_{mem}, T_{math})$.
- The longer of the two times demonstrates what limits performance:
 - If math time is longer, we say that a function is math limited;
 - If memory time is longer then it is memory limited.

Arithmetic Intensity

- How much time is spent in memory or math operations depends on both the algorithm and its implementation, as well as the processor's bandwidths.
- Memory time is equal to the number of bytes accessed in memory divided by the processor's memory bandwidth.
- Math time is equal to the number of operations divided by the processor's math bandwidth.

Arithmetic Intensity

- Thus, on a given processor a given algorithm is math limited if:
 - $T_{math} > T_{mem}$
 - $\frac{\#op}{BW_{math}} > \frac{\#bytes}{BW_{mem}}$
- By simple algebra, the inequality can be rearranged to:
 - $\frac{\#op}{\#bytes} > \frac{BW_{math}}{BW_{mem}}$
- The left-hand side: the algorithm's arithmetic intensity.
- The right-hand side: ops:byte ratio.

Arithmetic Intensity

- Arithmetic intensity: the ratio of algorithm implementation operations and the number of bytes accessed.
- Ops:byte ratio: the ratio of a processor's math and memory bandwidths.
- Thus, an algorithm is math limited on a given processor if the algorithm's arithmetic intensity is higher than the processor's ops:byte ratio.
- Conversely, an algorithm is memory limited if its arithmetic intensity is lower than the processor's ops:byte ratio.

Arithmetic Intensity

- Compare the algorithm's arithmetic intensity to the ops:byte ratio on an NVIDIA Volta V100 GPU.
 - V100 has a peak math rate of 125 FP16 Tensor TFLOPS;
 - An off-chip memory bandwidth of approx. 900 GB/s
 - An on-chip L2 bandwidth of 3.1 TB/s;
- So it has a ops:byte ratio between 40 and 139, depending on the source of an operation's data (on-chip or off-chip memory).

Operation	Arithmetic Intensity	limited by
Linear layer (4096 outputs, 1024 inputs, batch size 512)	315 FLOPS/B	arithmetic
Linear layer (4096 outputs, 1024 inputs, batch size 1)	1 FLOPS/B	memory
Max pooling with 3x3 window and unit stride	2.25 FLOPS/B	memory
ReLU activation	0.25 FLOPS/B	memory
Layer normalization	10 FLOPS/B	memory

Arithmetic Intensity

- Note that this type of analysis is a simplification, as we're counting only the algorithmic operations used.
- In practice, functions also contain instructions for operations not explicitly expressed in the algorithm, such as:
 - Memory access instructions;
 - Address calculation instructions;
 - Control flow instructions, and so on.

Limited by Lantency

- The arithmetic intensity and ops:byte ratio analysis assumes that a workload is sufficiently large to saturate a given processor's math and memory pipelines.
- However, if the workload is not large enough, or does not have sufficient parallelism, the processor will be under-utilized and performance will be limited by latency.
- For example:
 - Consider the launch of a single thread that will access 16 bytes and perform 16000 math operations.
 - While the arithmetic intensity is 1000 FLOPS/B and the execution should be math-limited on a V100 GPU, creating only a single thread grossly under-utilizes the GPU, leaving nearly all of its math pipelines and execution resources idle.

DNN Operation Categories

Elementwise Operations

- Elementwise operations may be unary or binary operations;
- The key is that layers in this category perform mathematical operations on each element independently of all other elements in the tensor.
- For example:
 - A ReLU layer returns $\max(0, x)$ for each x in the input tensor.
 - The element-wise addition of two tensors computes each output sum value independently of other sums.
- Layers in this category include most non-linearities (sigmoid, tanh, etc.), scale, bias, add, and others.
- These layers tend to be *memory-limited*, as they perform few operations per byte accessed.

Reduction Operations

- Reduction operations produce values computed over a range of input tensor values.
- For example:
 - Pooling layers compute values over some neighbourhoods in the input tensor.
 - Batch normalization computes the mean and standard deviation over a tensor before using them in operations for each output element.
 - SoftMax also falls into the reduction category.
- Typical reduction operations have a low arithmetic intensity and thus *are memory limited*.

Dot-Product Operations

- Operations in this category can be expressed as dot-products of elements from two tensors, usually a weight (learned parameter) tensor and an activation tensor.
- Examples:
 - Fully-connected layers are naturally expressed as matrix-vector and matrix-matrix multiplies.
 - Convolutions can also be expressed as collections of dot-products - one vector is the set of parameters for a given filter, and the other is an “unrolled” activation region to which that filter is being applied.
- Operations in the dot-product category can be math-limited if the corresponding matrices are large enough.
- However, for the smaller sizes, these operations end up being memory-limited. For example, a fully-connected layer applied to a single vector (a tensor for a mini-batch of size 1)) is memory limited.

Dot-Product Operations: Matrix Multiplication

- Compute $C = AB$ suppose:
 - A is an $M \times K$ matrix; (M rows and K columns)
 - B is an $K \times N$ matrix;
 - C is an $M \times N$ matrix;
- A total of $M \times N \times K$ fused multiply-adds (FMAs) are needed to compute the product. so a total of $2 \times M \times N \times K$ flops are required.
- The total number of byte scan in FP16: $2(M \times K + K \times N + M \times N)$
- Arithmetic intensity =
$$\frac{M \times N \times K}{(M \times K + K \times N + M \times N)}.$$

References

- <https://docs.nvidia.com/deeplearning/performance/dl-performance-gpu-background/index.html#understand-perf>
- <https://developer.nvidia.com/blog/nvidia-ampere-architecture-in-depth/>
- https://www.alcf.anl.gov/sites/default/files/2021-07/ALCF_A100_20210728%5B80%5D.pdf